

ICCS312 Algorithms and Tractability (Term 1/2021-22)

Lecture 5: Amortized Analysis


Sunsern Cheamanunkul

Mahidol University International College

September 28, 2021


- 1 Amortized Analysis
 - Aggregate method
 - Accounting method
 - Potential method
- 2 Example: Binary Counter
- 3 Discussion

Stack



- Stack Operations:
 - $\text{Push}(S, x)$ pushes object x onto stack S


Stack



- Stack Operations:

- $\text{Push}(S, x)$ pushes object x onto stack S
- $\text{Pop}(S)$ pops the top of (non-empty) stack S


Stack



- Stack Operations:

- $\text{Push}(S, x)$ pushes object x onto stack S
- $\text{Pop}(S)$ pops the top of (non-empty) stack S
- $\text{MultiPop}(S, k)$ pops the k top objects from S

Stack



- Stack Operations:
 - $\text{Push}(S, x)$ pushes object x onto stack S
 - $\text{Pop}(S)$ pops the top of (non-empty) stack S
 - $\text{MultiPop}(S, k)$ pops the k top objects from S
- What is the largest possible cost of a sequence of n stack operations?

Stack

- What is the largest possible cost of a sequence of n stack operations?

Stack

- What is the largest possible cost of a sequence of n stack operations?
- Simple worst-case bound analysis
 - The largest cost is $\Theta(n)$ from MultiPop
 - There are n stack operations
 - Total cost is $\Theta(n^2)$

Stack

- What is the largest possible cost of a sequence of n stack operations?
- Simple worst-case bound analysis
 - The largest cost is $\Theta(n)$ from MultiPop
 - There are n stack operations
 - Total cost is $\Theta(n^2)$
- Correct, but it's not tight (too pessimistic).

Stack

- What is the largest possible cost of a sequence of n stack operations?
- Simple worst-case bound analysis
 - The largest cost is $\Theta(n)$ from MultiPop
 - There are n stack operations
 - Total cost is $\Theta(n^2)$
- Correct, but it's not tight (too pessimistic).
- We can show that the **amortized** cost for each operation is just $\Theta(1)$.

Amortized Analysis

- Analyze a **sequence** of n operations
- Show that the *averaged* cost for each operation is small.

Amortized Analysis

- Analyze a **sequence** of n operations
- Show that the *averaged* cost for each operation is small.

Def.

The **amortized cost per operation** for a sequence of n operations is the total cost of the operations divided by n .

Amortized Analysis

- Analyze a **sequence** of n operations
- Show that the *averaged* cost for each operation is small.

Def.

The **amortized cost per operation** for a sequence of n operations is the total cost of the operations divided by n .

- Ex: If we have 100 operations at cost 1, followed by 1 operation at cost 100, the *amortized* cost per operation is $200/101 < 2$.

Amortized Analysis

- Three main techniques:
 1. The aggregate method – analyzes the total running time for a sequence of operations.
 2. The accounting method – charges extra and use to pay for future operations.
 3. The potential method – derives a *potential function* to describe the cost of the operations.

Stack with an extensible array

- Consider a stack implemented with an array.
- We have an array A with a variable top that points to the top of the stack.
- Operations:
 1. Push
 $A[top] = x;$
 $top++;$
 2. Pop
 $top--;$
 $\text{return } A[top];$

Stack with an extensible array

- When the array is full, we need to allocate a new larger array and copy the data over.
 - Sounds expensive!
- Some push operation will be more expensive than others.
- Luckily this doesn't happen too often.
- What can we say about the *amortized* cost for push?

Stack with an array

- Let's define our cost model first.

Cost model

- Write an item into the array costs 1.
- Read an item from the array costs 1.
- Cost for the growing array is the number of elements copied over.

Policy 1: Grow one by one

- When the array is full, increase its size by 1.

Policy 1: Grow one by one

- When the array is full, increase its size by 1.
- Consider n push operations. The total cost will be $1 + 2 + 3 + \dots + n = n(n + 1)/2$.

Policy 1: Grow one by one

- When the array is full, increase its size by 1.
- Consider n push operations. The total cost will be $1 + 2 + 3 + \dots + n = n(n + 1)/2$.
- The *amortized* cost per operation is $(n + 1)/2$ or $O(n)$.

Policy 2: Doubling

- When the array is full, double its size.

Policy 2: Doubling

- When the array is full, double its size.
- Consider n push operations.

Policy 2: Doubling

- When the array is full, double its size.
- Consider n push operations.
- Two types of costs:
 1. Total resize cost: $1 + 2 + 4 + 8 + \dots + 2^i$ for some $2^i < n$.
 - The sum is at most $2n - 1$.
 2. Total push cost is n from n pushes.
- Therefore, the total cost is $3n - 1 < 3n$.

Policy 2: Doubling

- When the array is full, double its size.
- Consider n push operations.
- Two types of costs:
 1. Total resize cost: $1 + 2 + 4 + 8 + \dots + 2^i$ for some $2^i < n$.
 - The sum is at most $2n - 1$.
 2. Total push cost is n from n pushes.
- Therefore, the total cost is $3n - 1 < 3n$.
- The amortized cost per operation is < 3 or $O(1)$.

Accounting method

- Idea: show that we have enough money pay for the total cost of n operations out of a bank account.
- Charge extra for *cheap* operations to pay for more *expensive* operations later.

Accounting method

- Idea: show that we have enough money pay for the total cost of n operations out of a bank account.
- Charge extra for *cheap* operations to pay for more *expensive* operations later.
- Let c_i be the true cost of the i -th operation.
- Show that by charging some cost c'_i for the i -th operation, we have

$$\sum_{1 \leq i \leq n} c_i \leq \sum_{1 \leq i \leq n} c'_i$$

- The *amortized* cost per operation is $(\sum_{1 \leq i \leq n} c'_i)/n$.


Accounting method

Claim.

With the doubling policy, if we charge \$3 for each push, we have enough money to pay for all the costs incurred during the n operations.

Proof.

For each push, we use \$1 to pay for the write right away and put \$2 into our account. Anytime, we need to double the array, from L to $2L$, we pay for it using the money in our bank account. We know we will always have enough money is because, after the last resizing, there were $L/2$ elements in the arrays. So there must be at least $L/2 * 2$ dollars in the account. ■



pay \$3

\$ 1 for push

\$ 2 saved in the account

since last resize,

\$ $(L/2 * 2)$ - to pay for resizing

$L \rightarrow 2L$

Accounting method

Claim.

With the doubling policy, if we charge \$3 for each push, we have enough money to pay for all the costs incurred during the n operations.

Proof.

For each push, we use \$1 to pay for the write right away and put \$2 into our account. Anytime, we need to double the array, from L to $2L$, we pay for it using the money in our bank account. We know we will always have enough money is because, after the last resizing, there were $L/2$ elements in the arrays. So there must be at least $L/2 * 2$ dollars in the account. ■

- Hence, the *amortized* cost per operation is 3.

Potential method

- Similar to the accounting method, but this looks at states of the data structure.
- Let h_i be the state of our data structure after the i -th operation and h_0 be the initial state.

Definition

Suppose we define a potential function Φ on states of a data structure such that $\Phi(h_0) = 0$ and $\Phi(h_i) \geq 0$ for all states h_i .

The *amortized* cost \tilde{c} of an operation is given by

$$\tilde{c} = c + (\Phi(h') - \Phi(h))$$

where c is the actual cost of the operation and h and h' is the states before and after the operation.

Stack with a fixed-size array

- Recall $\text{Push}(S,x)$, $\text{Pop}()$, and $\text{MultiPop}(S,x,k)$
- Propose: $\Phi(h_i)$ = number of items in the stack after the i -th operation.
- Observation:
 - $\Phi(h_0) = 0$ as initial stack is empty.
 - $\Phi(h_i) \geq 0$ as the number of items in the stack cannot be negative.

Stack with a fixed-size array

$\Phi(h_i)$ = number of items in the stack after the i -th operation.

- Push

- actual cost $c_i = 1$
- potential change: $\Phi(h_i) - \Phi(h_{i-1}) = 1$
- amortized cost: $\tilde{c} = c_i + (\Phi(h_i) - \Phi(h_{i-1})) = 1 + 1 = 2$

Stack with a fixed-size array

$\Phi(h_i)$ = number of items in the stack after the i -th operation.

- Pop
 - actual cost $c_i = 1$
 - potential change: $\Phi(h_i) - \Phi(h_{i-1}) = -1$
 - amortized cost: $\tilde{c} = c_i + (\Phi(h_i) - \Phi(h_{i-1})) = 1 - 1 = 0$



Stack with a fixed-size array

$\Phi(h_i)$ = number of items in the stack after the i -th operation.

- MultiPop

- actual cost $c_i = \min(k, |S|)$
- potential change: $\Phi(h_i) - \Phi(h_{i-1}) = -\min(k, |S|)$
- amortized cost: $\tilde{c} = c_i + (\Phi(h_i) - \Phi(h_{i-1})) = 0$

Stack with a fixed-size array



Example - Binary Counter

- Imagine we want to store a big binary counter in an array A
- Operation: Inc() – increment the counter by 1.
- Cost for flipping each bit is 1

count Val

	A[m]	A[m-1]	...	A[3]	A[2]	A[1]	A[0]	cost
0	0	0	...	0	0	0	0	
1	0	0	...	0	0	0	1	\$1
2	0	0	...	0	0	1	0	\$2
3	0	0	...	0	0	1	1	\$1
4	0	0	...	0	1	0	0	\$3
5	0	0	...	0	1	0	1	\$1
								\$2


Handwritten red annotations: A vertical line separates the 'count' and 'Val' columns. Red arrows indicate the sequence of bit flips for each increment: from row 0 to 1 (A[0]), from row 1 to 2 (A[0] and A[1]), from row 2 to 3 (A[0]), from row 3 to 4 (A[0], A[1], and A[2]), and from row 4 to 5 (A[0]).

Example - Binary Counter

- Simple worst-case bound:
 - $O(\log n)$ as each increment changes at most $\log n$
- Let's do amortized analysis of each increment.

Example - Binary Counter

Aggregate method



Counter Value	A[3]	A[2]	A[1]	A[0]	Total Cost
0	0	0	0	0	0
1	0	0	0	1	1
2	0	0	1	0	3
3	0	0	1	1	4
4	0	1	0	0	7
5	0	1	0	1	8
6	0	1	1	0	10
7	0	1	1	1	11

- How often do we flip each bit?

Example - Binary Counter

- $A[0]$ is flipped every time.
- $A[1]$ is flipped every other time.
- $A[2]$ is flipped every 4th time, and so on.

Example - Binary Counter

- $A[0]$ is flipped every time.
- $A[1]$ is flipped every other time.
- $A[2]$ is flipped every 4th time, and so on.
- So, after n increments,
 - The total number of $A[0]$ flips is n
 - The total number of $A[1]$ flips is $n/2$
 - The total number of $A[2]$ flips is $n/4$, and so on.

Example - Binary Counter

- $A[0]$ is flipped every time.
- $A[1]$ is flipped every other time.
- $A[2]$ is flipped every 4th time, and so on.
- So, after n increments,
 - The total number of $A[0]$ flips is n
 - The total number of $A[1]$ flips is $n/2$
 - The total number of $A[2]$ flips is $n/4$, and so on.
- Total cost is $n + n/2 + n/4 + \dots \leq 2n$.
- Thus, the *amortized* cost for each increment is ≤ 2 .

Example - Binary Counter

Accounting method

- Imagine each bit having its own bank account.
- Pay \$2 every time we flip $0 \rightarrow 1$.
 - \$1 for the actual cost and put \$1 into the account.
- When we flip $1 \rightarrow 0$, use the money from the account.
- We will always have enough money to pay for the increments.
- Hence, the amortized cost per increment is \$2.

* Each increment has exactly one $0 \rightarrow 1$ flip.

Example - Binary Counter

Potential method

$\Phi(h_i)$ = the number of 1-bits in the current count.

- Easy to check that:
 - $\Phi(h_0) = 0$
 - $\Phi(h_i) \geq 0$ for all h_i
- Proof: Leave this as exercise to you!

Discussion

- Suppose now it costs 2^k to flip the k -th bit
- In a sequence of n increments,
 - Simple worst-case bound gives $O(n)$ per increment.
 - What is the *amortized* cost per increment?

$2^2=4$	$2^1=2$	$2^0=1$	
0	0	0	\$1
0	0	1	\$3
0	1	0	\$1
0	1	1	\$7
1	0	0	

$$\text{total cost} \leq 2^0 \cdot n + 2^1 \frac{n}{2} + 2^2 \frac{n}{2^2} + \dots + 2^k \frac{n}{2^k}$$

where $k \leq \lceil \log n \rceil \leftarrow$ we perform n increments
only $\lceil \log n \rceil$ bits are affected.

$$\leq \underbrace{n + n + \dots + n}_{(k+1) \text{ term}}$$

$$= n \cdot (k+1) \leq n \cdot (\lceil \log n \rceil + 1)$$

So, Amortised cost per increment is

$$\frac{n \lceil \log n \rceil}{n} = O(\log n) \quad \square$$

bits affected

0 000

1 001

2 010

3 011

4 100

5 101

0

1

2

2

3

3